

System Programming Lab #8

2020-05-13

sp-tas

Lab Assignment #4 : Kernel Lab

- Download skeleton code & pdf from eTL
 - kernellab-handout.tar, kernellab-handout.pdf
- Hand In #1 – Setup (Done)
 - capture your development environment
 - Upload your capture image eTL
 - 압축파일 양식 : [학번]_이름]_kernellab_setup.tar (or .zip etc) **(including below files)**
 - filename for part #1 : [학번]_[이름]_kernellab_ptree.jpg (or .png, etc)
 - filename for part #2: [학번]_[이름]_kernellab_paddr.jpg (or .png, etc)
- Hand In #2 – Your Implementation
 - Upload your files eTL
 - 압축파일 양식 : [학번]_[이름]_kernellab.tar (or .zip, etc)
 - Ex) 2020-12345_홍길동_kernellab.tar
 - A zip file should include
 - (1) a tarball of your implementation directory (2) report
 - tarball 양식 : kernellab-[학번].tar.gz eg) kernellab-2020-12345.tar.gz
 - Report 양식 : [학번]_[이름]_kernellab_report.pdf (or .hwp, .txt etc)
- Please, **READ** the Hand-out and Lab material thoroughly!

Lab Assignment #4 : Kernel Lab

- Step 1. Setup (Done)
 - (part #0) Load my own kernel module
- Step 2. Implementation
 - (part #1) Tracing process tree from process id
 - (part #2) Finding physical address using virtual address
- Assigned : May 6
- Deadline for Step 1. Setup : May 13, 23:59:59 (Delay NOT allowed)
- Deadline for Step 2. Implementation : May 27, 23:59:59
- Delay policy : Same as before
- Lab sessions will be
 - 5/6: Kernel lab part #0, #1
 - 5/13: Kernel lab part #2 ← TODAY
 - 5/27 : Kernel lab Q&A session

Notice

- Please re-download the handout .tar and .pdf file
 - Makefile bug fixed
 - Source code and description for part #2 updated

Today's Lab

- Some useful tools in kernel programming
 - tmux, ctags, cscope
- Part #2. Finding physical address using virtual address
 - Remind translation of VA->PA in Computer Architecture
 - Assignment spec
 - How to begin
 - Testing your program
- Evaluation

tmux (terminal multiplexer)

- Installation
 - `sudo apt-get install tmux`
- Basic Commands

Command	Description
<code>tmux</code>	start tmux
<code>tmux new -s <name></code>	start tmux with <name>
<code>tmux ls</code>	shows the list of sessions
<code>tmux a #</code>	attach the detached-session
<code>tmux a -t <name></code>	attach the detached-session to <name>
<code>tmux kill-session -t <name></code>	kill the session <name>
<code>tmux kill-server</code>	kill the tmux server

Split windows with tmux

- **1. ctrl + b (or a)** to type tmux command
- 2. Split vertically : `%(shift 5)`
- 3. Split horizontally : `”(shift ‘)`
- <https://tmuxguide.readthedocs.io/en/latest/tmux/tmux.html>

Ctags

- What is Ctags?
 - A Tool that makes it easy to navigate big source code projects.
- Ctags generates **database** of tag file
 - for global variables, functions, macros, etc
 - to point where they are declared & defined
- Installation
 - `sudo apt-get install ctags (or exuberant-ctags)`
- Check
 - `ctags --version`
- Help
 - `Ctags --help`

Ctags – how to make tags file

- Steps

- 1. go to root directory of codes you want to navigate.
 - `cd /(where your root directory of code is)`
- 2. generate tags file
 - type `ctags -R` (recursive)
 - or `ctags file1, file2, ...`
- 3. Check tags file
 - `ls`

```
1. ta@sp3: ~/yschoi/malloclab/src (ssh)
ta@sp3:~/yschoi/malloclab/src$ ctags -R
ta@sp3:~/yschoi/malloclab/src$ ls
checkalign  clock.o  fcyc.o  ftimer.c  Makefile-handout  memlib.c  mm-explicit.c  mm.o  README-handout
checkalign.c  config.h  fsecs.c  ftimer.h  mdriver          memlib.h  mm.h           mm-test.c  tags
clock.c      fcyc.c   fsecs.h  ftimer.o  mdriver.c        memlib.o  mm-implicit.c  mm-tree.c
clock.h      fcyc.h   fsecs.o  Makefile  mdriver.o        mm.c      mm-naive.c     README
```

- 4. Remove tags file
 - `rm tags`

Ctags – how to use

- Case 1. In code file
 - 1. place cursor on the keyword you want to locate where it is defined
 - 2. type **ctrl +]**

```
227 if ((bp = mem_sbrk(size)) == (void *)-1)
228     return NULL;
229
230 /* Initialize free block header/footer and the epilogue header */
231 PUT(HDRP(bp), PACK(size, 0)); /* free block header */
232 PUT(FTRP(bp), PACK(size, 0)); /* free block footer */
233 PUT(HDRP(NEXT_BLKP(bp)), PACK(0, 1)); /* new epilogue header */
234
235 /* Coalesce if the previous block was free */
236 return coalesce(bp);
237 }
238 /* $end mmextendheap */
239
240 /*
241  * place - Place block of asize bytes at start of free block bp
242  *         and split if remainder would be at least minimum block size
243  */
244 /* $begin mmplace */
245 /* $begin mmplace-proto */
246 static void place(void *bp, size_t asize)
247 {
248     /* ... */
249 }
```

227:17 [62%]
"mm.c" 369L, 9640C

```
58 void *mem_sbrk(int incr)
59 {
60     char *old_brk = mem_brk;
61
62     if ( (incr < 0) || ((mem_brk + incr) > mem_max_addr)) {
63         errno = ENOMEM;
64         fprintf(stderr, "ERROR: mem_sbrk failed. Ran out of memory...\n");
65         return (void *)-1;
66     }
67     mem_brk += incr;
68     return (void *)old_brk;
69 }
70
71 /*
72  * mem_heap_lo - return address of the first heap byte
73  */
74 void *mem_heap_lo()
75 {
76     return (void *)mem_start_brk;
77 }
78
79 void *mem_heap_hi()
80 {
81     return (void *)mem_end_brk;
82 }
```

58:1 [61%]
'memlib.c' 101L, 2270C

- 3. type **ctrl + t** to go back

Ctags – how to use

- Case 2. In tags file
 - 1. **vi tags**
 - 2. type **:tj [tag name]** to find

```
35 FTRP mm-implicit.c 71;" d file:
36 FTRP mm.c 71;" d file:
37 GET mm-implicit.c 62;" d file:
38 GET mm.c 62;" d file:
1:1 [Top]
:tj mem_sbrk
```

- 3. type **:po** to comeback

```
58 void *mem_sbrk(int incr)
59 {
60     char *old_brk = mem_brk;
61
62     if ( (incr < 0) || ((mem_brk + incr) > mem_max_addr)) {
63         errno = ENOMEM;
64         fprintf(stderr, "ERROR: mem_sbrk failed. Ran out of memory...\n");
65         return (void *)-1;
66     }
67     mem_brk += incr;
68     return (void *)old_brk;
69 }
70
71 /*
72  * mem_heap_lo - return address of the first heap byte
73  */
74 void *mem_heap_lo()
75 {
76     return (void *)mem_start_brk;
77 }
58:1 [61%]
:po
```

Cscope

- A tool to navigate in big source code.
- Diff with Ctags?
 - Able to locate **functions where they are called** too.
- Installation
 - `sudo apt-get install cscope`
- Check
 - `cscope --version`
- Help
 - `cscope --help`

```
ta@sp3:~/yschoi/malloclab/src$ cscope --version
cscope: version 15.8b
ta@sp3:~/yschoi/malloclab/src$
```

Cscope – how to make cscope database file

- Steps
 - 1. go to root directory of codes you want to navigate.
 - `cd /(where your root directory of code is)`
 - 2. generate cscope database file
 - `find ./ -name '*[cCsShH]' > file_list`
 - `cscope -i file_list`
 - 3. Check cscope.out file
 - `ls`

```
ta@sp3:~/yschoi/malloclab/src$ find ./ -name '*[cCsShH]' > file_list
ta@sp3:~/yschoi/malloclab/src$ cscope -i file_list
ta@sp3:~/yschoi/malloclab/src$ ls
checkalign  cscope.out  fsecs.h    Makefile-handout  memlib.o    mm.o
checkalign.c  fcyc.c      fsecs.o    mdriver            mm.c         mm-test.c
clock.c       fcyc.h      ftimer.c   mdriver.c          mm-explicit.c mm-tree.c
clock.h       fcyc.o      ftimer.h   mdriver.o          mm.h         README
clock.o       file_list   ftimer.o   memlib.c           mm-implicit.c README-handout
config.h      fsecs.c     Makefile   memlib.h           mm-naive.c
```

- 4. Remove tags file
 - `rm cscope.out file_list`

Cscope – how to use

- 1. type **cscope** to execute
- 2. type **ctrl+d** to break out

```
Cscope version 15.8b                                Press the ? key for help

Find this C symbol:
Find this global definition:
Find functions called by this function:
Find functions calling this function: mem_sbrk
Find this text string:
Change this text string:
Find this egrep pattern:
Find this file:
Find files #including this file:
Find assignments to this symbol:
```

```
Functions calling this function: mem_sbrk

File      Function      Line
mm-explicit.c requestMoreSpace 172 ptrNewBlock = (void *)((unsigned int )mem_sbrk(totalSize
- 4);
1 mm-explicit.c mm_init      204 if (mem_sbrk(initsize) == (void *)-1) {
2 mm-implicit.c mm_init      99 if ((heap_listp = mem_sbrk(4*WSIZE)) == NULL)
3 mm-implicit.c extend_heap  227 if ((bp = mem_sbrk(size)) == (void *)-1)
4 mm-naive.c   mm_malloc      62 void *p = mem_sbrk(newsize);
5 mm-test.c   mm_init      43 mem_sbrk(64000);
6 mm-test.c   mm_malloc      54 void *p = mem_sbrk(newsize);
7 mm-tree.c   mm_init      726 if (mem_sbrk(HEAP_INITSIZE) == NULL)
8 mm-tree.c   mm_realloc    826 if (mem_sbrk(grow_size) == NULL)
9 mm-tree.c   mm_malloc      875 if (mem_sbrk(block_size) == NULL)

Find this C symbol:
Find this global definition:
Find functions called by this function:
Find functions calling this function:
Find this text string:
Change this text string:
Find this egrep pattern:
Find this file:
Find files #including this file:
Find assignments to this symbol:
```

Downloading Linux kernel source code

- <https://mirrors.edge.kernel.org/pub/linux/kernel/v4.x/linux-4.15.tar.gz>
- You can download the source code using wget
- Usage: `$ wget <source>`

```
ubuntu@sp:~/src$ wget https://mirrors.edge.kernel.org/pub/linux/kernel/v4.x/linux-4.15.tar.gz
--2020-05-12 14:59:30-- https://mirrors.edge.kernel.org/pub/linux/kernel/v4.x/linux-4.15.tar.gz
Resolving mirrors.edge.kernel.org (mirrors.edge.kernel.org)... 147.75.95.133, 2604:1380:3000:1500::1
Connecting to mirrors.edge.kernel.org (mirrors.edge.kernel.org)|147.75.95.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 157656459 (150M) [application/x-gzip]
Saving to: 'linux-4.15.tar.gz'

linux-4.15.tar.gz                19%[=====>] 28.93M  2.53MB/s  eta 50s
```

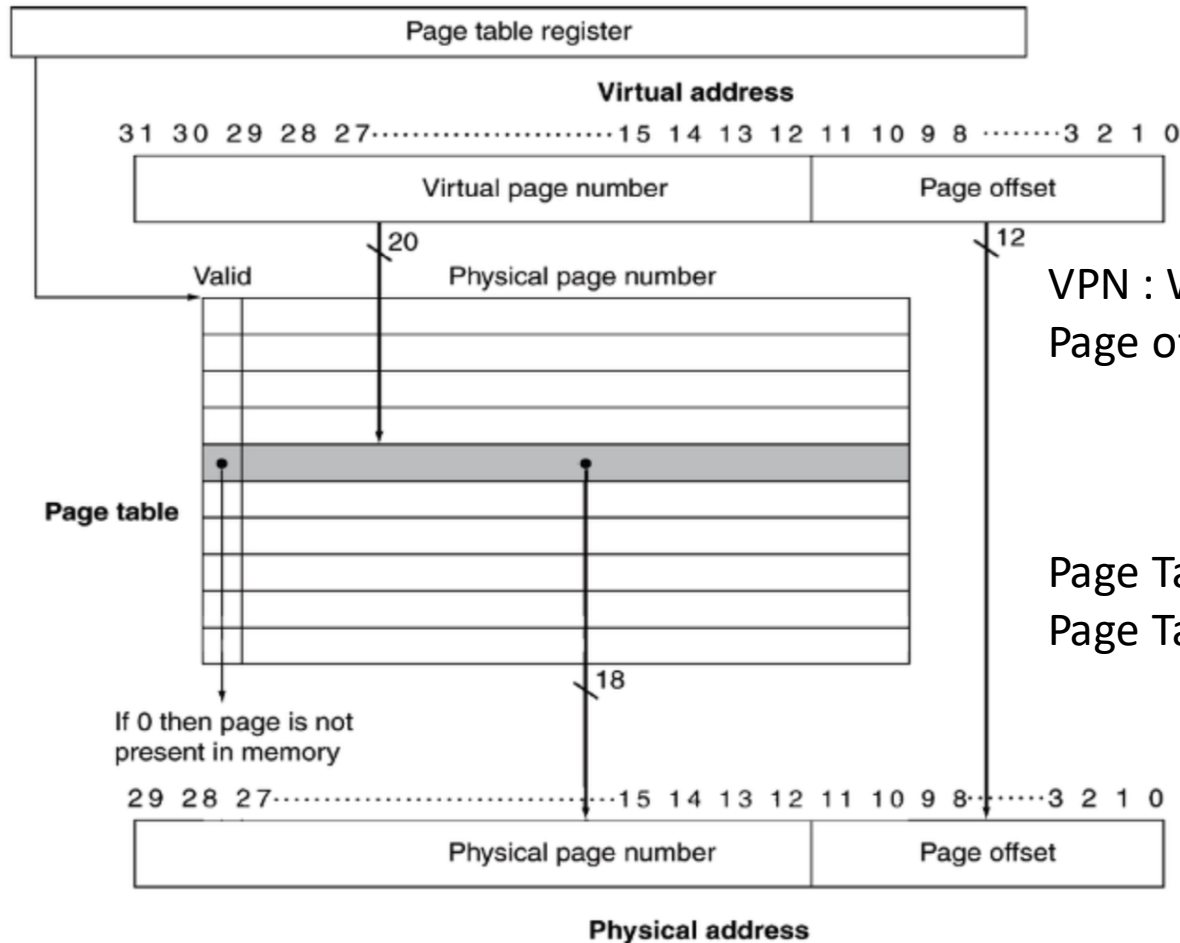
Playing with Linux kernel source code

- Uncompress the tar archive
 - `$ tar -xzf linux-4.15.tar.gz`
- Go to the root of the source code
 - `$ cd linux-4.15/`
- Create tag files using Linux kernel's makefile
 - `$ make tags cscope ARCH=x86_64 -j2`
- Now we are all set!
 - Example: `$ vi -t debugfs_create_file`

Playing with Linux kernel source code

- `$ vi -t <symbol>`
- Inside vim
 - `g + ctrl +]`
 - `:tj`
 - `ctrl + t`
- `$ cscope`

Virtual to Physical Address Translation



VPN : Which page to look for


Page offset : Where in the page to look for

Page Table : Contains process page info

Page Table Entry : Record of Page Table

Example

- Virtual to Physical Translation
 - 4KB page size, 20bit virtual address
 - VA : 0xF13B4



Multilevel Page Tables

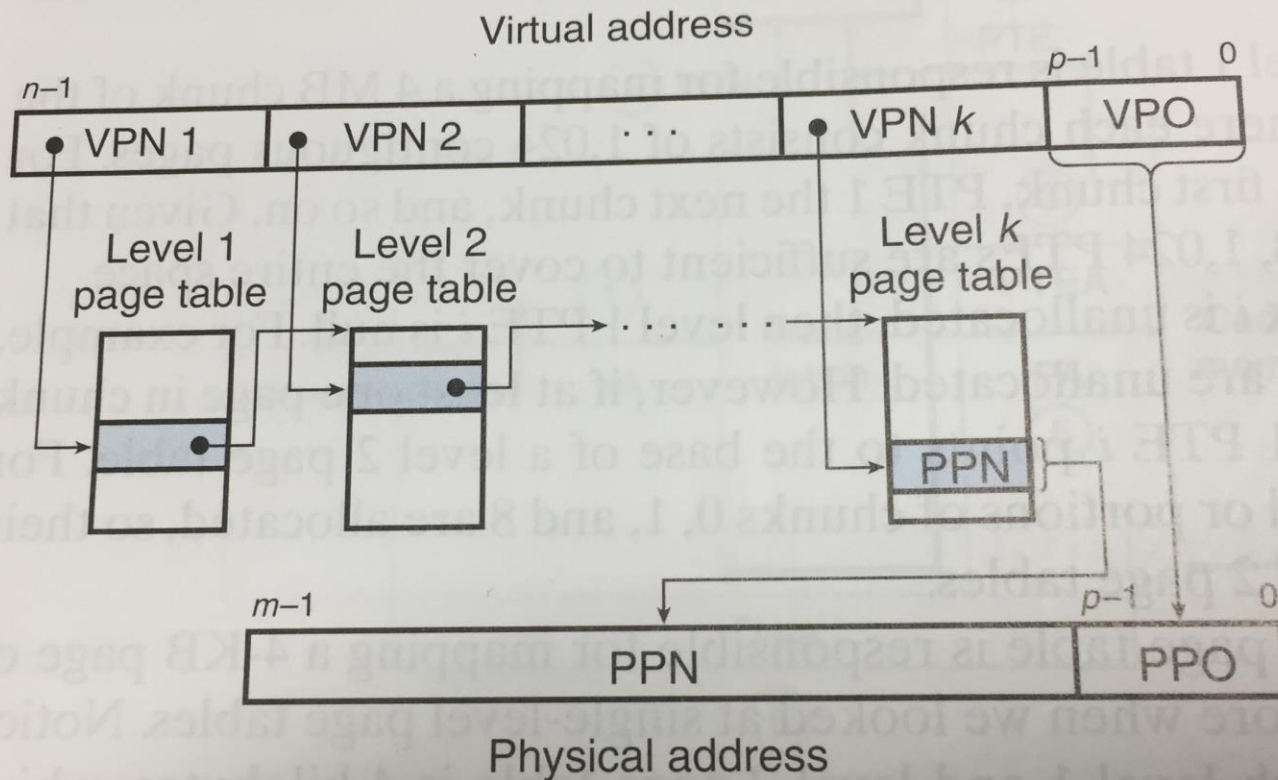


Figure 9.18 Address translation with a k -level page table.

Part2. Finding physical address using virtual address

- Spec
 - In app.c
 - makes a virtual address mapped to predefined physical address
 - In your kernel module
 - 1. get pid of app and virtual address
 - 2. returns physical address
 - In app.c
 - Compares return value from kernel module with predefined value.

Hints

- Page walk API
 - <Linux source root>/arch/x86/include/asm/pgtable.h
- Look for the schemes how virtual address is translated to physical address
- Page walk procedure in Linux 4.15.0
 - pgd -> p4d -> pud -> pmd -> pte

Testing your program

- Step
 - 0. `sudo su`
 - 1. `make`
 - 2. `./app`

```
root@yschoi-VirtualBox:/home/yschoi/kernellab_full/solution/paddr# sudo su
root@yschoi-VirtualBox:/home/yschoi/kernellab_full/solution/paddr# make
make -C /lib/modules/4.15.0-47-generic/build M=/home/yschoi/kernellab_full/solution/paddr modules;
make[1]: Entering directory '/usr/src/linux-headers-4.15.0-47-generic'
  CC [M]  /home/yschoi/kernellab_full/solution/paddr/dbfs_paddr.o
  Building modules, stage 2.
  MODPOST 1 modules
  CC      /home/yschoi/kernellab_full/solution/paddr/dbfs_paddr.mod.o
  LD [M]  /home/yschoi/kernellab_full/solution/paddr/dbfs_paddr.ko
make[1]: Leaving directory '/usr/src/linux-headers-4.15.0-47-generic'
gcc -o app app.c;
sudo insmod dbfs_paddr.ko
root@yschoi-VirtualBox:/home/yschoi/kernellab_full/solution/paddr# ./app
vaddr { 7ffbf301000 } paddr { 0 }
vaddr { 7ffbf301000 } paddr { 0 }
vaddr { 7ffbf301000 } paddr { 234512000 }
[TEST CASE]      PASS
root@yschoi-VirtualBox:/home/yschoi/kernellab_full/solution/paddr#
```

Asked Questions

- 1. ptree
 - (1) init process의 이름이 systemd
 - 상관 없습니다.
 - (2) Corner cases?
 - 채점 기준: 다음 명령이 올바르게 작동 하는지 확인, 모든 pid는 valid
 - # make
 - # cd /sys/kernel/debug/ptree
 - # echo pid1 >> input
 - # cat ptree
 - # echo pid2 >> input
 - # cat ptree
 - ...
 - # echo pidn >> input
 - # cat ptree
 - # cd -
 - # make clean

References

- Tmux guide
 - <https://tmuxguide.readthedocs.io/en/latest/tmux/tmux.html>
- Ctags
 - <https://bowbowbow.tistory.com/15>
- Cscope
 - <https://harryp.tistory.com/131>
- Address Translation, Multilevel Page table
 - P.849~855, R. E. Briant, D. R. O'Hallaron, Computer Systems, A programmer's perspective 3rd edition