

System Programming Lab #10

2020-05-27

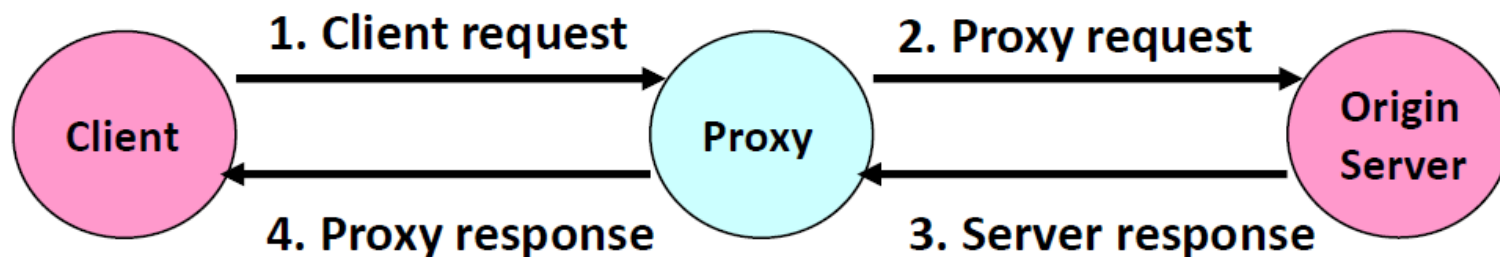
sp-tas

Lab Assignment #5 : Proxy Lab

- Download skeleton code & pdf from eTL
proxylab-handout.tar, proxylab-handout.pdf
- Hand In
 - First change **STUNO** to yours defined in Makefile
 - 'make handin' command will generate a tarball automatically
 - 구현 디렉토리 압축파일: 학번-proxylab.tar eg) 2020-12345-proxylab.tar
 - Upload your files eTL
 - 압축파일 양식 : [학번]_[이름]_proxylab.tar (or .zip, etc)
 - Ex) 2020-12345_홍길동_proxylab.tar
 - A zip file should include
 - (1) a tarball of your implementation directory (2) report
 - tarball 양식 : [학번]-proxylab.tar eg) 2020-12345-proxylab.tar
 - Report 양식 : [학번]_[이름]_proxylab_report.pdf (or .doc, .txt etc)
- Please, **READ** the Hand-out and Lab material thoroughly!
- Assigned : May 27
- Deadline : June 17, 23:59:00 (NO Delay Allowed)

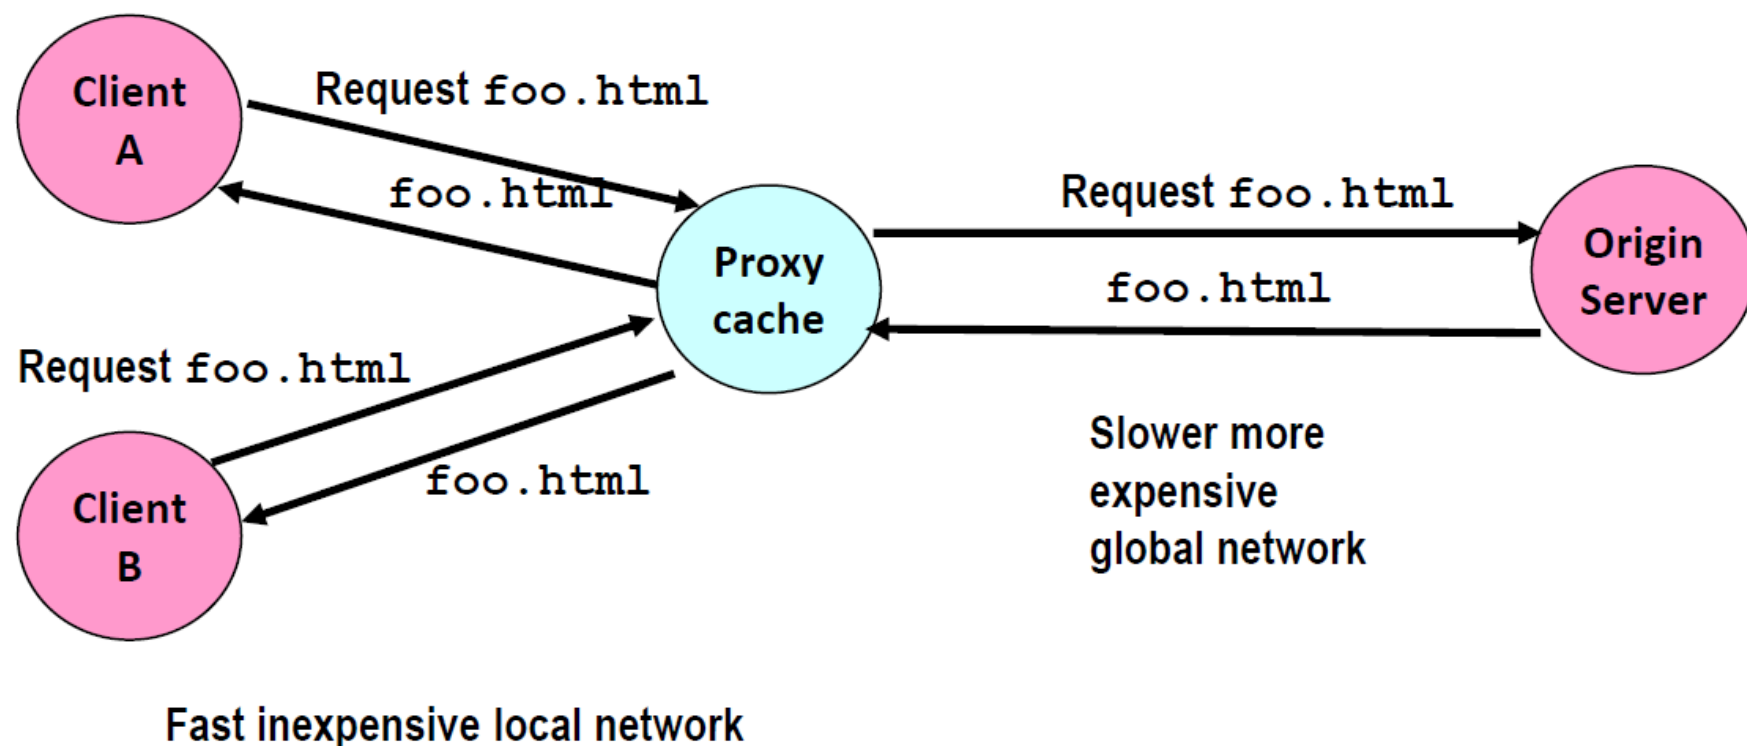
Proxies

- A **proxy** is an intermediary between a client and an **origin server**
 - To the client, the proxy acts like a server
 - To the server, the proxy acts like a client



Why Proxies?

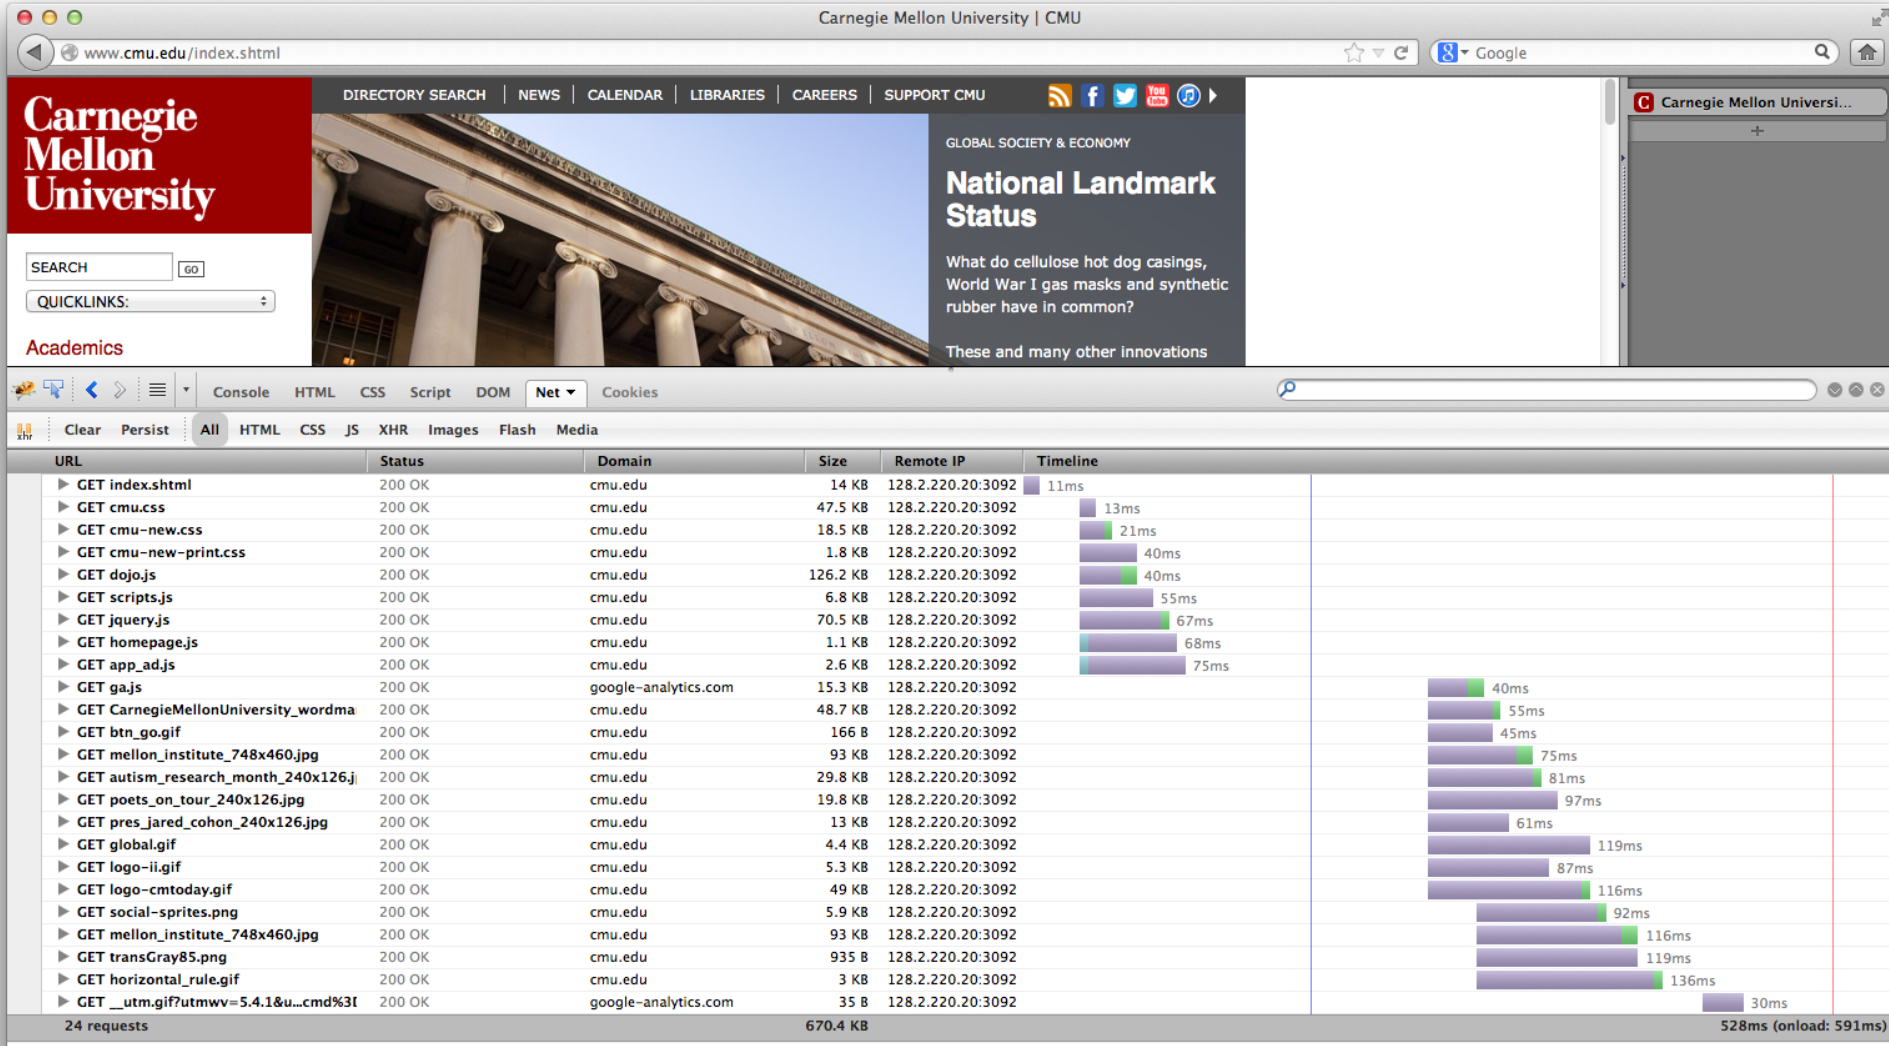
- Can perform useful functions as requests and responses pass by
 - Examples: Caching, logging, anonymization, filtering, transcoding



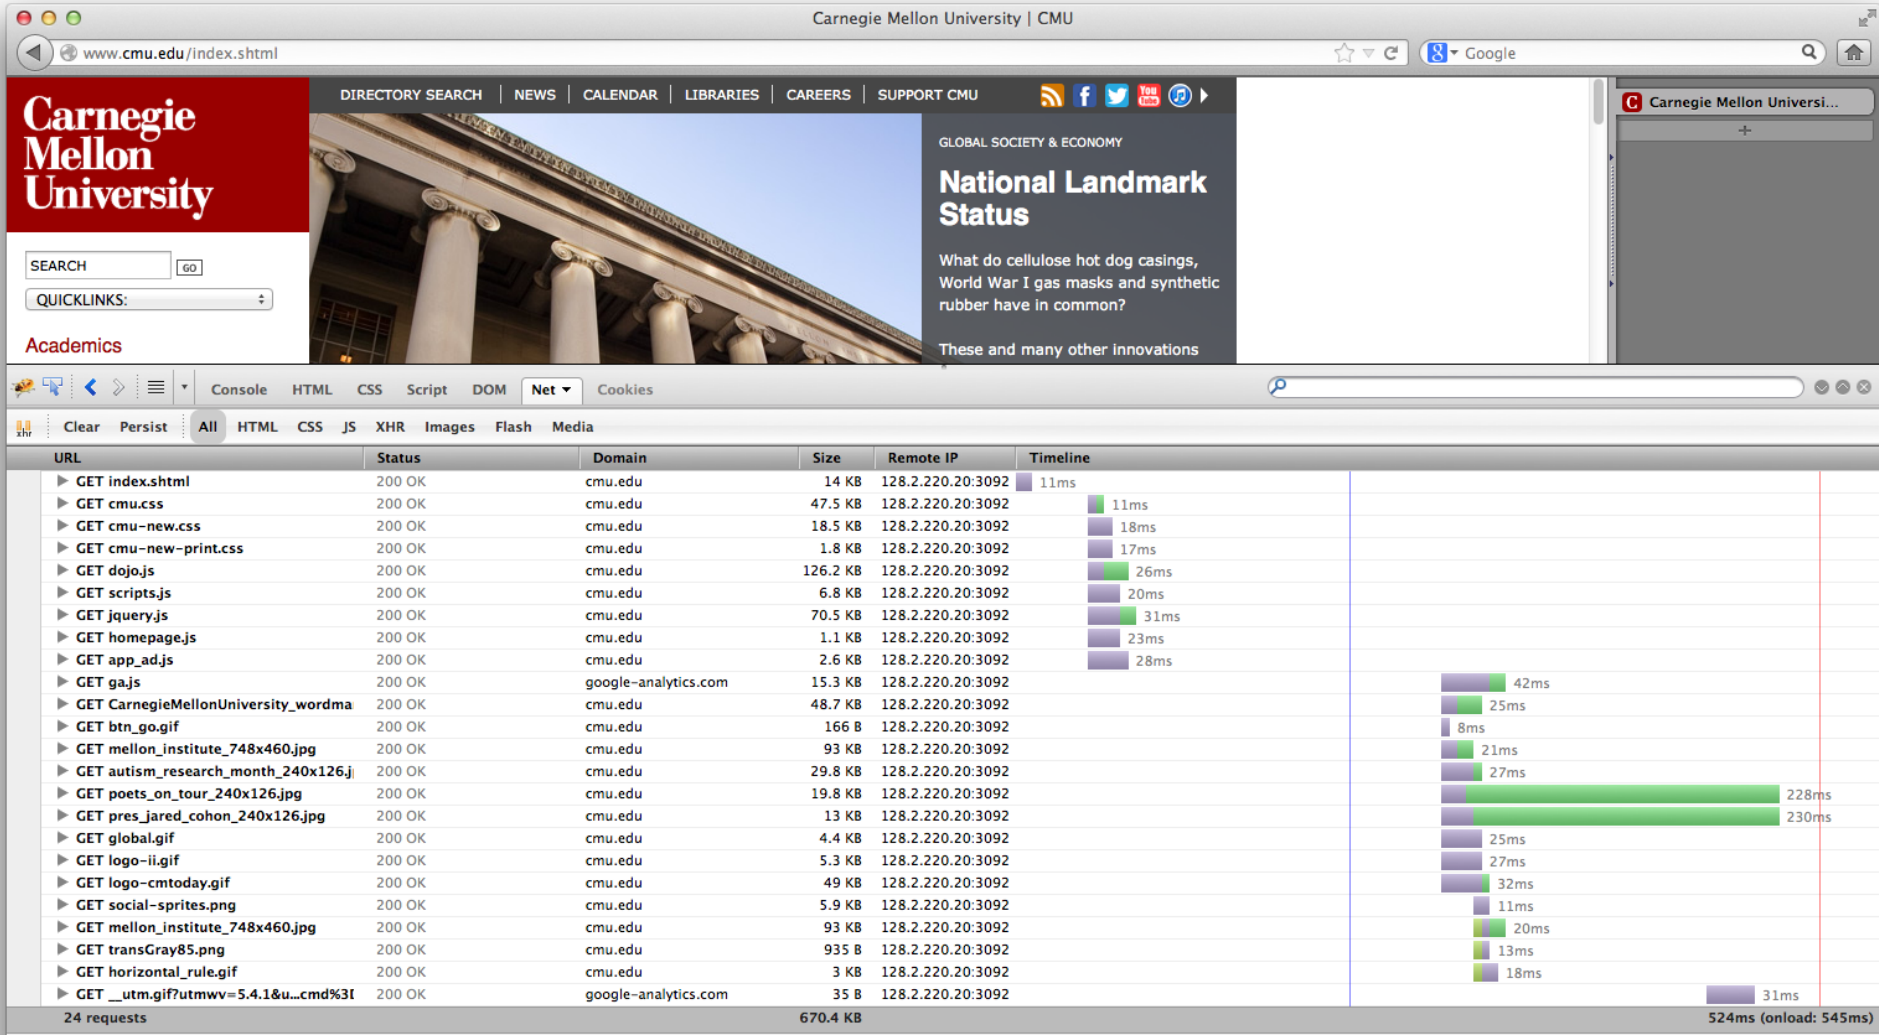
How the Web Really Works

- In reality, a single HTML page today may depend on 10s or 100s of support files (images, stylesheets, scripts, etc.)
- Builds a good argument for concurrent servers
 - Just to load a single modern webpage, the client would have to wait for 10s of back-to-back request
 - I/O is likely slower than processing, so back
- Caching is simpler if done in pieces rather than whole page
 - If only part of the page changes, no need to fetch old parts again
 - Each object (image, stylesheet, script) already has a unique URL that can be used as a key

Sequential Proxy



Concurrent Proxy



You will implement

- Write a simple HTTP proxy that caches web objects
- Part 1: Implementing a sequential web Proxy
 - **Basic HTTP operation & socket programming**
 - set up the proxy to accept incoming connections
 - read and parse requests
 - forward requests to web servers
 - read the servers' responses
 - forward those responses to the corresponding clients
- Part 2: Dealing with multiple concurrent requests
 - upgrade your proxy to deal with multiple **concurrent** connections
 - multi-threading
- Part 3: Caching web objects
 - add caching to your proxy using a simple main memory cache of recently accessed web content
 - cache individual objects, not the whole page
 - **Use an LRU eviction policy**
 - your caching system must allow for concurrent reads while maintaining consistency

Guide to start your implementation

- `int main(int argc, char *argv[])`
 - initialize everything such as data structure
 - checking port number
 - establish listening requests
 - when a client connects, spawn a new thread to handle it

```
1  #include <stdio.h>
2
3  /* Recommended max cache and object sizes */
4  #define MAX_CACHE_SIZE 1049000
5  #define MAX_OBJECT_SIZE 102400
6
7  /* You won't lose style points for including this long line in your code */
8  static const char *user_agent_hdr = "User-Agent: Mozilla/5.0 (X11; Linux x86_64; rv:10.0.3) Gecko/20120305 Firefox/10.0.3\r\n";
9
10 int main()
11 {
12     printf("%s", user_agent_hdr);
13     return 0;
14 }
15
```

Guide to start your implementation

- TAs implemented following structures and functions

```
typedef struct {  
    |  
} Request;  
  
void *handle_client(void *vargp);  
void initialize_struct(Request *req);  
void parse_request(char request[MAXLINE], Request *req);  
void parse_absolute(Request *req);  
void parse_relative(Request *req);  
void parse_header(char header[MAXLINE], Request *req);  
void assemble_request(Request *req, char *request);  
int get_from_cache(Request *req, int clientfd);  
void get_from_server(Request *req, char request[MAXLINE], int clientfd, rio_t rio_to_client);  
void close_wrapper(int fd);  
void print_full(char *string);  
void print_struct(Request *req);  
  
typedef struct CachedItem CachedItem;  
  
struct CachedItem {  
    |  
};  
  
typedef struct {  
    |  
} CacheList;  
  
extern void cache_init(CacheList *list);  
extern void cache_URL(char *URL, void *item, size_t size, CacheList *list);  
extern void evict(CacheList *list);  
extern CachedItem *find(char *URL, CacheList *list);  
extern void move_to_front(char *URL, CacheList *list);  
extern void print_URLs(CacheList *list);  
extern void cache_destruct(CacheList *list);
```

Use csapp.[ch] functions

- Also, csapp.[ch] codes are included! yeah!

```
/* Sockets interface wrappers */
int Socket(int domain, int type, int protocol);
void Setsockopt(int s, int level, int optname, const void *optval, int optlen);
void Bind(int sockfd, struct sockaddr *my_addr, int addrlen);
void Listen(int s, int backlog);
int Accept(int s, struct sockaddr *addr, socklen_t *addrlen);
void Connect(int sockfd, struct sockaddr *serv_addr, int addrlen);
```

```
/* Protocol independent wrappers */
void Getaddrinfo(const char *node, const char *service,
                 const struct addrinfo *hints, struct addrinfo **res);
void Getnameinfo(const struct sockaddr *sa, socklen_t salen, char *host,
                 size_t hostlen, char *serv, size_t servlen, int flags);
void Freeaddrinfo(struct addrinfo *res);
void Inet_ntop(int af, const void *src, char *dst, socklen_t size);
void Inet_pton(int af, const char *src, void *dst);
```

```
/* DNS wrappers */
struct hostent *Gethostbyname(const char *name);
struct hostent *Gethostbyaddr(const char *addr, int len, int type);
```

```
/* Pthreads thread control wrappers */
void Pthread_create(pthread_t *tidp, pthread_attr_t *attrp,
                   void * (*routine)(void *), void *argp);
void Pthread_join(pthread_t tid, void **thread_return);
void Pthread_cancel(pthread_t tid);
void Pthread_detach(pthread_t tid);
void Pthread_exit(void *retval);
pthread_t Pthread_self(void);
void Pthread_once(pthread_once_t *once_control, void (*init_function)());
```

```
/* Rio (Robust I/O) package */
ssize_t rio_readn(int fd, void *usrbuf, size_t n);
ssize_t rio_writen(int fd, void *usrbuf, size_t n);
void rio_readinitb(rio_t *rp, int fd);
ssize_t rio_readnb(rio_t *rp, void *usrbuf, size_t n);
ssize_t rio_readlineb(rio_t *rp, void *usrbuf, size_t maxlen);
```

```
/* Wrappers for Rio package */
ssize_t Rio_readn(int fd, void *usrbuf, size_t n);
void Rio_writen(int fd, void *usrbuf, size_t n);
void Rio_readinitb(rio_t *rp, int fd);
ssize_t Rio_readnb(rio_t *rp, void *usrbuf, size_t n);
ssize_t Rio_readlineb(rio_t *rp, void *usrbuf, size_t maxlen);
```

```
/* Reentrant protocol-independent client/server helpers */
int open_clientfd(char *hostname, char *port);
int open_listenfd(char *port);
```

```
/* Wrappers for reentrant protocol-independent client/server helpers */
int Open_clientfd(char *hostname, char *port);
int Open_listenfd(char *port);
```

Checking Your Work

- Auto grader
 - `./driver.sh` will run the tests:
 - Ability to pull basic web pages from a server
 - Handle a (concurrent) request while another request is still pending
 - Fetch a web page again from your cache after the server has been stopped
 - This should help answer the question “is this what my proxy is supposed to do?”
 - Please don’t use this grader to definitively test your proxy; there are many things not tested here

Checking Your Work

- Test your proxy liberally
 - The web is full of special cases that want to break your proxy
 - Generate a port for yourself with `./port-for-user.pl [sp ID]`
 - Generate more ports for web servers and such with `./free-port.sh`
- Create a handin file with *make handin*
 - First you should change STUNO defined in Makefile to your student number
 - Will create a tar file for you with the contents of your proxylab-handin folder

Telnet/cURL Demo

- Telnet
 - Interactive remote shell – like ssh without security
 - Must build HTTP request manually
 - This can be useful if you want to test response to malformed headers

Double-Enter →

```
ubuntu@sp:~$ telnet apache.org 80
Trying 40.79.78.1...
Connected to apache.org.
Escape character is '^]'.
GET /index.html HTTP/1.0

HTTP/1.1 200 OK
Date: Wed, 27 May 2020 06:59:34 GMT
Server: Apache/2.4.18 (Ubuntu)
Last-Modified: Wed, 27 May 2020 06:10:30 GMT
ETag: "14a3e-5a69b11aa7bb0"
Accept-Ranges: bytes
Content-Length: 84542
Vary: Accept-Encoding
Cache-Control: max-age=3600
Expires: Wed, 27 May 2020 07:59:34 GMT
Connection: close
Content-Type: text/html

<!DOCTYPE html>
<html lang="en">
<head>
```

Telnet/cURL Demo

- cURL
 - “URL transfer library” with a command line program
 - Builds valid HTTP requests for you!

```
ubuntu@sp:~$ curl http://apache.org
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="utf-8">
  <meta http-equiv="X-UA-Compatible" content="IE=edge">
  <meta name="viewport" content="width=device-width, initial-scale=1">
  <meta name="description" content="Home page of The Apache Software Foundation">

  <link rel="apple-touch-icon" sizes="57x57" href="/favicons/apple-touch-icon-57x57.png">
  <link rel="apple-touch-icon" sizes="60x60" href="/favicons/apple-touch-icon-60x60.png">
  <link rel="apple-touch-icon" sizes="72x72" href="/favicons/apple-touch-icon-72x72.png">
```

- Can also be used to generate HTTP proxy requests:

```
ubuntu@sp:~$ curl --proxy localhost:52184 http://apache.org > curl.txt
```

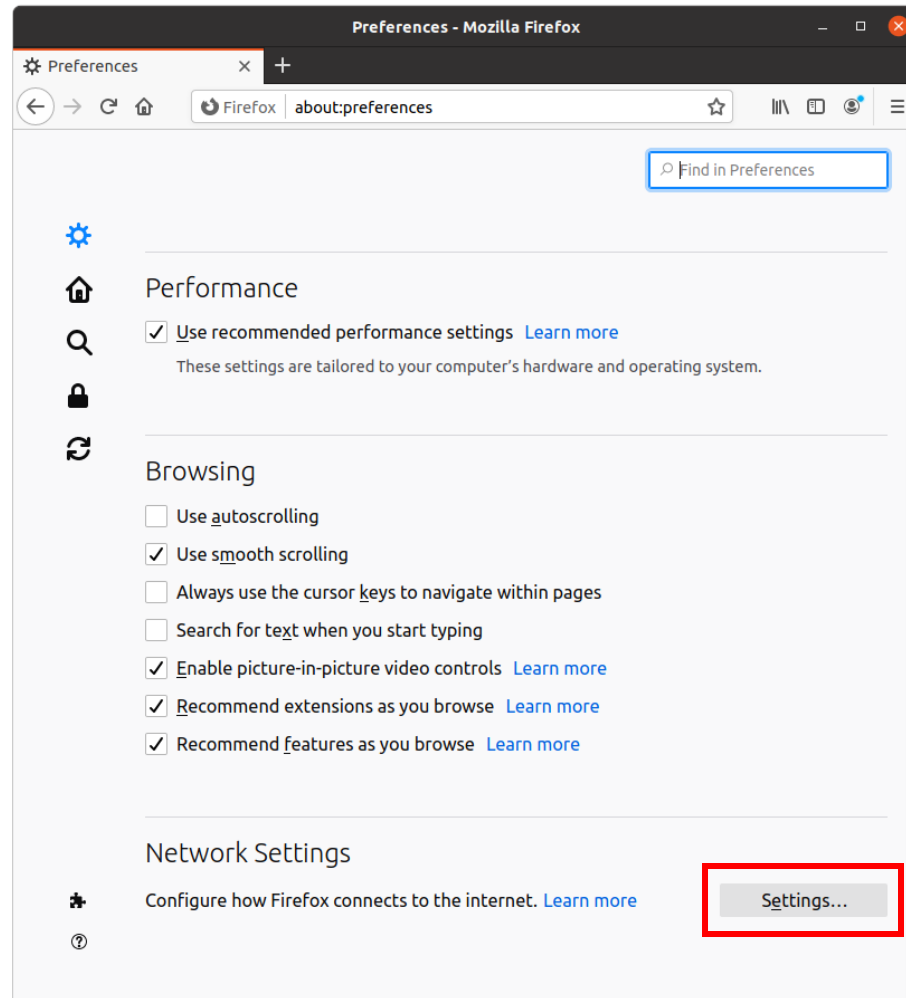
% Total	% Received	% Xferd	Average Speed	Dload	Upload	Time Total	Time Spent	Time Left	Current Speed
100	16962	100	16962	0	0	21.9M	0	--:--:--	16.1M

```
ubuntu@sp:~$ head -n 3 curl.txt
```

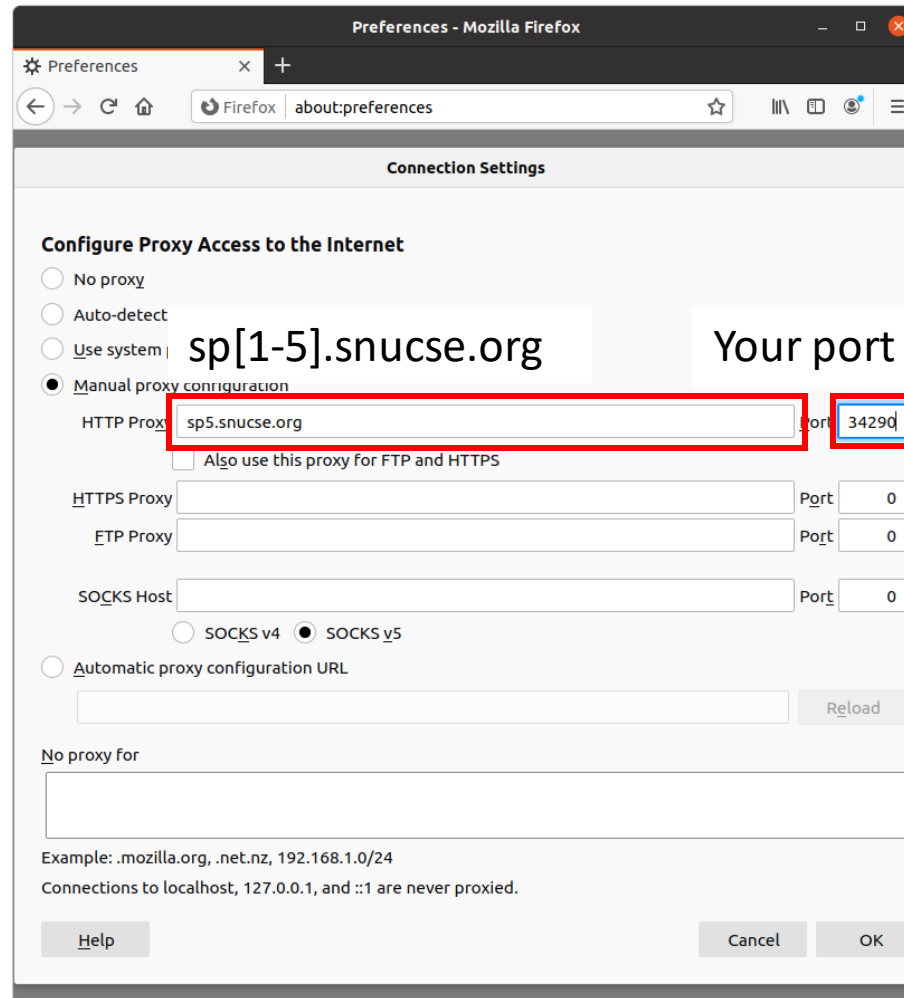
```
Ie6).mR.GÄDU(PU"$f^>b>3%+
                                A|h|_NeeEE+e@Iee&*SevY)eeee?w"Q"xoo/zPI}u
xK_:":<a7^oK$]<Q"#%%a~)S)%=baaeTeeWdeeeeD7veeeee-e"Σ^eeefez~e~l}---Vzozzv7ee獍'W}s-
qOm_4&&Wexofzf3damTeee*aaSy7#C#[2Deee5e_C9ee]ellweeePecqqgy*e5xe={=6eT>UpK++ehDekOee2ee
```

Gzipped (encoded) data from proxy

Testing with Web Browser (Firefox)



Testing with Web Browser (Firefox)



Test manually using curl

- Manually testing following real pages
 - <http://neverssl.com>
 - <http://example.com>
 - <http://apache.org>
 - <http://gnu.org>
- You should always use **`./port-for-user.pl username`** when testing your proxy manually

Evaluation

- Total Score: 80 points
- Basic Correctness (40 points)
 - basic proxy operation (auto graded)
- Concurrency (15 points)
 - handling concurrent requests (auto graded)
- Cache (15 points)
 - working cache (auto graded)
- Report (10 points)
 - describes the goal of proxy lab and how to implement for each part
 - what you learn in this lab
 - what was difficult
 - what was surprising and so on