

# System Programming Lab #1

---

2020-03-25

sp-tas





# Case Study: Library Interpositioning

- **Library interpositioning** : powerful linking technique that allows programmers to intercept calls to arbitrary functions
- **Interpositioning can occur at:**
  - Compile time: When the source code is compiled
  - Link time: When the relocatable object files are statically linked to form an executable object file
  - Load/run time: When an executable object file is loaded into memory, dynamically linked, and then executed.



# Some Interpositioning Applications

## ■ Security

- Confinement (sandboxing)
  - Interpose calls to libc functions.
- Behind the scenes encryption
  - Automatically encrypt otherwise unencrypted network connections.

## ■ Monitoring and Profiling

- Count number of calls to functions
- Characterize call sites and arguments to functions
- Malloc tracing
  - Detecting memory leaks
  - **Generating address traces**



# Example program

```
#include <stdio.h>
#include <stdlib.h>
#include <malloc.h>

int main()
{
    free(malloc(10));
    printf("hello, world\n");
    exit(0);
}
```

hello.c

- **Goal:** trace the addresses and sizes of the allocated and freed blocks, without modifying the source code.
- **Three solutions:** interpose on the `lib malloc` and `free` functions at compile time, link time, and load/run time.



# Compile-time Interpositioning

```
#ifdef COMPILETIME
/* Compile-time interposition of malloc and free using C
 * preprocessor. A local malloc.h file defines malloc (free)
 * as wrappers mymalloc (myfree) respectively.
 */

#include <stdio.h>
#include <malloc.h>

/*
 * mymalloc - malloc wrapper function
 */
void *mymalloc(size_t size, char *file, int line)
{
    void *ptr = malloc(size);
    printf("%s:%d: malloc(%d)=%p\n", file, line, (int)size,
ptr);
    return ptr;
}
```

# Compile-time Interpositioning

```
#define malloc(size) mymalloc(size, __FILE__, __LINE__ )
#define free(ptr) myfree(ptr, __FILE__, __LINE__ )

void *mymalloc(size_t size, char *file, int line);
void myfree(void *ptr, char *file, int line);
```

`malloc.h`

```
linux> make helloc
gcc -O2 -Wall -DCOMPILETIME -c mymalloc.c
gcc -O2 -Wall -I. -o helloc hello.c mymalloc.o
linux> make runc
./helloc
hello.c:7: malloc(10)=0x501010
hello.c:7: free(0x501010)
hello, world
```



# Link-time Interpositioning

```
#ifdef LINKTIME
/* Link-time interposition of malloc and free using the
static linker's (ld) "--wrap symbol" flag. */

#include <stdio.h>

void *__real_malloc(size_t size);
void __real_free(void *ptr);

/*
 * __wrap_malloc - malloc wrapper function
 */
void *__wrap_malloc(size_t size)
{
    void *ptr = __real_malloc(size);
    printf("malloc(%d) = %p\n", (int)size, ptr);
    return ptr;
}
```



# Link-time Interpositioning

```
linux> make hello1
gcc -O2 -Wall -DLINKTIME -c mymalloc.c
gcc -O2 -Wall -Wl,--wrap,malloc -Wl,--wrap,free \
-o hello1 hello.c mymalloc.o
linux> make run1
./hello1
malloc(10) = 0x501010
free(0x501010)
hello, world
```

- The “-Wl” flag passes argument to linker
- Telling linker “--wrap,malloc” tells it to resolve references in a special way:
  - Refs to `malloc` should be resolved as `__wrap_malloc`
  - Refs to `__real_malloc` should be resolved as `malloc`



```
#ifdef RUNTIME
/* Run-time interposition of malloc and free based on
 * dynamic linker's (ld-linux.so) LD_PRELOAD mechanism */
#define _GNU_SOURCE
#include <stdio.h>
#include <stdlib.h>
#include <dlfcn.h>

void *malloc(size_t size)
{
    void *(*mallocp)(size_t size);
    char *error;
    void *ptr;

    /* get address of libc malloc */
    if (!mallocp) {
        mallocp = dlsym(RTLD_NEXT, "malloc");
        if ((error = dlerror()) != NULL) {
            fputs(error, stderr);
            exit(1);
        }
    }
    ptr = mallocp(size);
    fprintf(stderr, "malloc(%d) = %p\n", (int)size, ptr);
    return ptr;
}
```

## Load/Run-time Interpositioning

mymalloc.c



# Load/Run-time Interpositioning

```
linux> make hellor
gcc -O2 -Wall -DRUNTIME -shared -fPIC -o mymalloc.so mymalloc.c
gcc -O2 -Wall -o hellor hello.c
linux> make runr
(LD_PRELOAD="/usr/lib/x86_64-linux-gnu/libdl.so ./mymalloc.so"
./hellor)
malloc(10) = 0x559a34eca260
free(0x559a34eca260)
hello, world
```

- The `LD_PRELOAD` environment variable tells the dynamic linker to resolve unresolved refs (e.g., to `malloc`) by looking in `libdl.so` and `mymalloc.so` first.
  - `libdl.so` necessary to resolve references to the `dlopen` functions.
- `-shared`: telling linker to make output as a shared objective (`.so`)
- `-fPIC`: Position-Independent Code



# Interpositioning Recap

## ■ Compile Time

- Apparent calls to malloc/free get macro-expanded into calls to mymalloc/myfree

## ■ Link Time

- Use linker trick to have special name resolutions
  - malloc → \_\_wrap\_malloc
  - \_\_real\_malloc → malloc

## ■ Load/Run Time

- Implement custom version of malloc/free that use dynamic linking to load library malloc/free under different names





# Dynamic Memory Management

**`void *malloc(size_t size)`**

`malloc` allocates `size` bytes of memory on the process' heap and returns a pointer to it that can subsequently be used by the process to hold up to `size` bytes. The contents of the memory are undefined.

**`void *calloc(size_t nmemb, size_t size)`**

`calloc` allocates `nmemb*size` bytes of memory on the process' heap and returns a pointer to it that can subsequently be used by the process to hold up to `size` bytes. The contents of the memory are set to zero.

**`void *realloc(void *ptr, size_t size)`**

`realloc` changes the size of the memory block pointed to by `ptr` to `size` bytes. The contents are copied up to `min(size, old size)`, the rest is undefined.

**`void free(void *ptr)`**

`free` explicitly frees a previously allocated block of memory.

# Dynamic Memory Management

```
#include <stdlib.h>

void main(void) {
    void *p;
    char *str;
    int *A;

    // allocated 1024 bytes of memory
    p = malloc(1024);

    // allocated an integer array with 500 integer
    A = (int*)calloc(500, sizeof(int));

    // allocate a string with 16 characters...
    str = (char*)malloc(16*sizeof(char));

    // ...then resize that string to hold 512 characters
    str = (char*)realloc(str, 512*sizeof(char));

    // finally, free all allocated memory
    free(p);
    free(A);
    free(str);
}
```

example1.c

# Shared library loading interfaces

- `#include <dlfcn.h>`
- `void *dlopen(const char *filename, int flags)`
  - `dlopen` loads the dynamic shared object (shared library) file named by the null-terminated string `filename` and returns an opaque “handle” for the loaded object.
  - **RTLD\_LAZY** – perform lazy binding
  - **RTLD\_NOW** – all undefined symbols in the shared object are resolved before `dlopen` returns
- `void *dlsym(void *handle, const char *symbol)`
  - `dlsym` takes a “handle” of a dynamic loaded shared object returned by `dlopen` and returns the address where that symbol is loaded into memory
  - **RTLD\_DEFAULT** – find the first occurrence of the desired symbol
  - **RTLD\_NEXT** – find the next occurrence of the desired symbol in the search order after the current object
- `int dlclose(void *handle)`
  - `dlclose` decrements the reference count on the dynamically loaded shared object referred to by `handle`



```

#ifdef RUNTIME
/* Run-time interposition of malloc and free based on
 * dynamic linker's (ld-linux.so) LD_PRELOAD mechanism */
#define _GNU_SOURCE
#include <stdio.h>
#include <stdlib.h>
#include <dlfcn.h>

void *malloc(size_t size)
{
    void *(*mallocp)(size_t size);
    char *error;
    void *ptr;

    /* get address of libc malloc */
    if (!mallocp) {
        mallocp = dlsym(RTLD_NEXT, "malloc");
        if ((error = dlerror()) != NULL) {
            fputs(error, stderr);
            exit(1);
        }
    }
    ptr = mallocp(size);
    fprintf(stderr, "malloc(%d) = %p\n", (int)size, ptr);
    return ptr;
}

```

## Load/Run-time Interpositioning





# Lab Assignment #1 – Linker Lab

- Download skeleton code from eTL
- Hand In
  - Upload your files **eTL**
  - A tarball of your implementation (20/20/20/+10 pts for each part)
  - A report (10 pts)
- PLEASE, **READ** the Hand-out!!!
- Assigned: 3. 25
- Deadline: 4. 8, 11:59:59 PM
- Lab #2 (4/1) will be Q&A session

Shared library를 통해  
Malloc, calloc, realloc, free를 intercept하여  
기존의 함수 작업을 하면서  
memory tracking도 수행

# (Part 1) Tracing dynamic memory allocation

test1.c

```
#include <stdlib.h>

void main(void) {
    void *a;

    a = malloc(1024);
    a = malloc(32);
    free(malloc(1));
    free(a);
}
```

output

```
stuXXX@spN ~/linklab/part1 $ make run test1
[0001] Memory tracer started.
[0002]          (nil) : malloc( 1024 ) = 0xb87010
[0003]          (nil) : malloc(  32 ) = 0xb87420
[0004]          (nil) : malloc(   1 ) = 0xb87450
[0005]          (nil) : free( 0xb87450 )
[0006]          (nil) : free( 0xb87420 )
[0007]
[0008] Statistics
[0009]   allocated_total      1057
[0010]   allocated_avg        352
[0011]   freed_total         0
[0012]
[0013] Memory tracer stopped.
stuXXX@spN ~/linklab/part1 $
```

## (Part 2) Tracing unfreed memory

test1.c

```
#include <stdlib.h>
```

```
void main(void) {  
    void *a;
```

output

```
    a = malloc(1024);  
    a = malloc(32);  
    free(malloc(1));  
    free(a);  
}
```

```
stuXXX@spN ~/linklab/part2 $ make run test1  
[0001] Memory tracer started.  
[0002]          (nil) : malloc( 1024 ) = 0x2415060  
[0003]          (nil) : malloc( 32 ) = 0x24154c0  
[0004]          (nil) : malloc( 1 ) = 0x2415540  
[0005]          (nil) : free( 0x2415540 )  
[0006]          (nil) : free( 0x24154c0 )  
[0007]  
[0008] Statistics  
[0009]   allocated_total      1057  
[0010]   allocated_avg       352  
[0011]   freed_total         33  
[0012]  
[0013] Non-deallocated memory blocks  
[0014]   block              size      ref cnt   caller  
[0015]   0x2415060          1024        1      ??? : 0  
[0016]  
[0017] Memory tracer stopped.  
stuXXX@spN ~/linklab/part2 $
```

## (Part 3) Pinpointing call locations

test1.c

```
#include <stdlib.h>
```

```
void main(void) {  
    void *a;
```

```
    a = malloc(1024);  
    a = malloc(32);  
    free(malloc(1));  
    free(a);  
}
```

stuXXX@spN ~/linklab/part3 \$ make run test1

[0001] Memory tracer started.

[0002] **main:6** : malloc( 1024 ) = 0x14f0060

[0003] **main:10** : malloc( 32 ) = 0x14f04c0

[0004] **main:1d** : malloc( 1 ) = 0x14f0540

[0005] **main:25** : free( 0x14f0540 )

[0006] **main:2d** : free( 0x14f04c0 )

[0007]

[0008] Statistics

[0009] allocated\_total 1057

[0010] allocated\_avg 352

[0011] freed\_total 33

[0012]

[0013] Non-deallocated memory blocks

| [0014] | block     | size | ref cnt | caller        |
|--------|-----------|------|---------|---------------|
| [0015] | 0x14f0060 | 1024 | 1       | <b>main:6</b> |

[0016]

[0017] Memory tracer stopped.

stuXXX@spN ~/linklab/part3 \$

# (Part 3) Pinpointing call locations

```
~/linklab/part3 $ make run test1
[0001] Memory tracer started.
[0002] main:6 : malloc( 1024 ) = 0x14f0060
[0003] main:10 : malloc( 32 ) = 0x14f04c0
[0004] main:1d : malloc( 1 ) = 0x14f0540
[0005] main:25 : free( 0x14f0540 )
[0006] main:2d : free( 0x14f04c0 )
[0007]
[0008] Statistics
[0009]   allocated_total      1057
[0010]   allocated_avg        352
[0011]   freed_total          33
[0012]
[0013] Non-deallocated memory blocks
[0014]   block                size      ref cnt    caller
[0015]   0x14f0060            1024        1      main:6
[0016]
[0017] Memory tracer stopped.
~/linklab/part3 $
```

# (Part 3) Pinpointing call locations

```
~/linklab/test $ objdump -d test1
```

```
...
```

```
00000000004004a0 <main>:
```

|                |                       |              |                                  |
|----------------|-----------------------|--------------|----------------------------------|
| 4004a0:        | 53                    | push         | %rbx                             |
| 4004a1:        | bf 00 04 00 00        | mov          | \$0x400,%edi                     |
| <b>4004a6:</b> | <b>e8 e5 ff ff ff</b> | <b>callq</b> | <b>400490 &lt;malloc@plt&gt;</b> |
| 4004ab:        | bf 20 00 00 00        | mov          | \$0x20,%edi                      |
| <b>4004b0:</b> | <b>e8 db ff ff ff</b> | <b>callq</b> | <b>400490 &lt;malloc@plt&gt;</b> |
| 4004b5:        | bf 01 00 00 00        | mov          | \$0x1,%edi                       |
| 4004ba:        | 48 89 c3              | mov          | %rax,%rbx                        |
| <b>4004bd:</b> | <b>e8 ce ff ff ff</b> | <b>callq</b> | <b>400490 &lt;malloc@plt&gt;</b> |
| 4004c2:        | 48 89 c7              | mov          | %rax,%rdi                        |
| <b>4004c5:</b> | <b>e8 96 ff ff ff</b> | <b>callq</b> | <b>400460 &lt;free@plt&gt;</b>   |
| 4004ca:        | 48 89 df              | mov          | %rbx,%rdi                        |
| <b>4004cd:</b> | <b>e8 8e ff ff ff</b> | <b>callq</b> | <b>400460 &lt;free@plt&gt;</b>   |
| 4004d2:        | 31 c0                 | xor          | %eax,%eax                        |
| 4004d4:        | 5b                    | pop          | %rbx                             |
| 4004d5:        | c3                    | retq         |                                  |
| 4004d6:        | 66 2e 0f 1f 84 00 00  | nopw         | %cs:0x0(%rax,%rax,1)             |
| 4004dd:        | 00 00 00              |              |                                  |

```
~/linklab/test $
```



# (Part 3) Pinpointing call locations

```
~/linklab/part3 $ make run test1
[0001] Memory tracer started.
[0002] main:6 : malloc( 1024 ) = 0x14f0060
[0003] main:10 : malloc( 32 ) = 0x14f04c0
[0004] main:1d : malloc( 1 ) = 0x14f0540
[0005] main:25 : free( 0x14f0540 )
[0006] main:2d : free( 0x14f04c0 )
[0007]
[0008] Statistics
[0009]   allocated_total      1057
[0010]   allocated_avg        352
[0011]   freed_total          33
[0012]
[0013] Non-deallocated memory blocks
[0014]   block                size      ref cnt    caller
[0015]   0x14f0060            1024        1      main:6
[0016]
[0017] Memory tracer stopped.
~/linklab/part3 $
```

## (Part 3) Pinpointing call locations

```
//  
// return the PC of the callsite to the dynamic memory management function  
//  
//  fname      pointer to character array to hold function name  
//  fnlen      length of character array  
//  ofs        pointer to offset to hold PC offset into function  
//  
// returns  
//  0          on success  
//  <0        on error  
//  
int get_callinfo(char *fname, size_t fnlen, unsigned long long *ofs);  
  
                                ~/linklab/part3/callinfo.h
```

```
int get_callinfo(char *fname, size_t fnlen,  
                unsigned long long *ofs)  
{  
    return -1;  
}  
  
                                ~/linklab/part3/callinfo.c
```

get\_callinfo() 를 작성하여 구현

# (Part 3) Backtracing library

실습 서버에 설치됨

- `#include <libunwind.h>`
  - <https://www.nongnu.org/libunwind/docs.html>
  - <https://github.com/libunwind/libunwind.git>
- `int unw_get_proc_name(unw_cursor_t *cp, char *bufp, size_t len, unw_word_t *offp);`
  - Get name of current procedure
  - Returns the name of the procedure that created the stack frame identified by argument cp
- `int unw_getcontext(unw_context_t *ucp);`
  - Get current machine-state
- `int unw_init_local(unw_cursor_t *c, unw_context_t *ctxt);`
  - Initialize cursor for local unwinding
- `int unw_step(unw_cursor_t *cp);`
  - Step to next (older) frame

# (Part 3) Backtracing library

```
#define UNW_LOCAL_ONLY
#include <libunwind.h>

static void print_backtrace(void)
{
    unw_context_t context;
    unw_cursor_t cursor;
    unw_word_t off, ip, sp;
    unw_proc_info_t pip;
    char procname[256];
    int ret;

    if (unw_getcontext(&context))
        return;

    if (unw_init_local(&cursor, &context))
        return;

    //???
}

void ccc(void)
{
    print_backtrace();
}

void bbb()
{
    ccc();
}

void aaa()
{
    bbb();
}

int main(void)
{
    aaa();
    return 0;
}
```

output

```
vi aaa.c
user100@SystemProgramming:~$ gcc aaa.c -lunwind
user100@SystemProgramming:~$ ./a.out
print_backtrace
ccc
bbb
aaa
main
__libc_start_main
_start
user100@SystemProgramming:~$
```

libunwind를 통해  
함수 호출 스택을 따라갈 수 있고  
함수 정보를 확인

## (Part 3) Pinpointing call locations

test1.c

```
#include <stdlib.h>
```

```
void main(void) {  
    void *a;
```

```
    a = malloc(1024);  
    a = malloc(32);  
    free(malloc(1));  
    free(a);  
}
```

stuXXX@spN ~/linklab/part3 \$ make run test1

[0001] Memory tracer started.

[0002] **main:6** : malloc( 1024 ) = 0x14f0060

[0003] **main:10** : malloc( 32 ) = 0x14f04c0

[0004] **main:1d** : malloc( 1 ) = 0x14f0540

[0005] **main:25** : free( 0x14f0540 )

[0006] **main:2d** : free( 0x14f04c0 )

[0007]

[0008] Statistics

[0009] allocated\_total 1057

[0010] allocated\_avg 352

[0011] freed\_total 33

[0012]

[0013] Non-deallocated memory blocks

| [0014] | block     | size | ref cnt | caller        |
|--------|-----------|------|---------|---------------|
| [0015] | 0x14f0060 | 1024 | 1       | <b>main:6</b> |

[0016]

[0017] Memory tracer stopped.

stuXXX@spN ~/linklab/part3 \$

# (Bonus) Detect and ignore illegal deallocations

- Detect double- free / illegal free

test4.c - test case for bonus part

```
#include <stdlib.h>

void main(void) {
    void *a;

    a = malloc(1024);
    free(a);
    free(a);
    free((void*)0x1706e90);
}
```

output

```
stuXXX@spN ~/linklab/bonus $ make run test4
[0001] Memory tracer started.
[0002]      main:6   : malloc( 1024 ) = 0x1b30060
[0003]      main:11  : free( 0x1b30060 )
[0004]      main:19  : free( 0x1b30060 )
[0005]      *** DOUBLE_FREE *** (ignoring)
[0006]      main:23  : free( 0x1706e90 )
[0007]      *** ILLEGAL_FREE *** (ignoring)
[0008]
[0009] Statistics
[0010]   allocated_total      1024
[0011]   allocated_avg       1024
[0012]   freed_total         1024
[0013]
[0014] Memory tracer stopped.
stuXXX@spN ~/linklab/bonus $
```

Free / Realloc 할 때

# Utilities

- Read someone else's code

- `memlog.c`

- `int mlog(int pc, const char* fmt, ...)`

- `memlist.c`

- `item *new_list(void)`
  - `void free_list(void)`
  - `item *alloc(item *list, void *ptr, size_t size)`
  - `item *dealloc(item *list, void *ptr)`
  - `item *find(item *list, void *ptr)`
  - `void dump_list(item *list)`

# Q&A

- 과제 설명 파일 먼저 읽기
- 과제 압축 풀고 Test 폴더에서 make 먼저
- Printf error
  - fprintf or mlog 사용
- test폴더의 textx는 사용하지 않음



# (Part 1) Tracing dynamic memory allocation

test1.c

```
#include <stdlib.h>

void main(void) {
    void *a;

    a = malloc(1024);
    a = malloc(32);
    free(malloc(1));
    free(a);
}
```

output

```
stuXXX@spN ~/linklab/part1 $ make run test1
[0001] Memory tracer started.
[0002]          (nil) : malloc( 1024 ) = 0xb87010
[0003]          (nil) : malloc( 32 ) = 0xb87420
[0004]          (nil) : malloc( 1 ) = 0xb87450
[0005]          (nil) : free( 0xb87450 )
[0006]          (nil) : free( 0xb87420 )
[0007]
[0008] Statistics
[0009]   allocated total      1057
[0010]   allocated_avg       352
[0011]   freed_total        0
[0012]
[0013] Memory tracer stopped.
stuXXX@spN ~/linklab/part1 $
```

Realloc의 경우, realloc size를 allocated\_total에 더함

## (Part 3) Pinpointing call locations

test1.c

```
#include <stdlib.h>
```

```
void main(void) {  
    void *a;
```

```
    a = malloc(1024);  
    a = malloc(32);  
    free(malloc(1));  
    free(a);  
}
```

stuXXX@spN ~/linklab/part3 \$ make run test1

[0001] Memory tracer started.

[0002] **main:6** : malloc( 1024 ) = 0x14f0060

[0003] **main:10** : malloc( 32 ) = 0x14f04c0

[0004] **main:1d** : malloc( 1 ) = 0x14f0540

[0005] **main:25** : free( 0x14f0540 )

[0006] **main:2d** : free( 0x14f04c0 )

[0007]

[0008] Statistics

[0009] allocated\_total 1057

[0010] allocated\_avg 352

[0011] freed\_total 33

[0012]

[0013] Non-deallocated memory blocks

| [0014] | block     | size | ref cnt | caller        |
|--------|-----------|------|---------|---------------|
| [0015] | 0x14f0060 | 1024 | 1       | <b>main:6</b> |

[0016]

Test.c에서 alloc, dealloc 함수는 main에서 호출된다고 가정

# Skeleton code snippet

```
//  
// init - this function is called once when the shared library is loaded  
//  
__attribute__((constructor))  
void init(void)  
{  
    char *error;  
  
    LOG_START();  
  
    // initialize a new list to keep track of all memory (de-)allocations  
    // (not needed for part 1)  
    list = new_list();  
  
    // ...  
}  
  
//  
// fini - this function is called once when the shared library is unloaded  
//  
__attribute__((destructor))  
void fini(void)  
{  
    // ...  
  
    LOG_STATISTICS(OL, OL, OL);  
  
    LOG_STOP();  
  
    // free list (not needed for part 1)  
    free_list(list);  
}
```

# 다음 시간에

- Linklab Q&A?
- 과제 기한 : 4월 8일, 11:59:59 PM