

System Programming

Linker Lab: Memtrace

Introduction

In this lab, we use the at-load-time library interpositioning to trace calls to the dynamic memory management system of a compiled binary program. Your memory tracer will be able to log all dynamic memory management calls of a program, and identify memory blocks that a program does not free. We then use a low-level libraries to identify the call sites that allocate the unfreed memory.

If you are still not satisfied, you can augment your memory tracer to prevent incorrect memory deallocations of random binaries in an optional bonus question.

By solving this lab, you will learn a lot about library interpositioning and simple management of data within a shared library. A lot of the code is provided already, this lab will also force you (and teach you how) to read someone else's code. We will also make use of a debugger library that will show you how to pinpoint exact program points in your running programs.

The lab may require a significant effort if you are not familiar with C programming. We provide enough hints and support in the lab sessions so that everyone can solve the lab with a good mark. Nevertheless, start early and ask early!

Logistics

Installation

Download the tarball from

```
eTL : linklab.tgz
```

and unpack it into a directory of your choice. In the :qples below we assume the directory name is linklab:

```
devel@gentoo ~ $ mkdir linklab
devel@gentoo ~ $ cd linklab
devel@gentoo ~/linklab $ tar xvzf linklab.tgz
devel@gentoo ~/linklab $ ls -l
```

Compiling, Running and Testing

Makefiles in the various directories assist you with submitting your compiling, running, and testing your implementations. Do not modify the Makefiles unless explicitly instructed. Run `make help` to find out which commands the Makefile support.

A set of test programs is provided with which you can test your solution. Note that we will use a different test set to evaluate your submission..

```
~/linklab/part1 $ make help
```

Submission

To submit your solution run `make submit` in the top directory. You must fill in your student ID and name in the file `STUDENT.ID`

```
~/linklab $ make submit
```

Also, you should submit a report about your implementation. You should include description about how to implement for each part of the lab, what was difficult, and wat was surprising, and so on. The report is 10 points.

Upload a zip file of your submission by eTL. The file should include a tarball of your implementation and a report.

Dynamic Memory Management

Overview

In many languages, memory is explicitly allocated and sometimes even deallocated by the programmer. The POSIX standard defines the following four dynamic memory management functions that we want to intercept.

`void *malloc(size_t size)`

`malloc` allocates `size` bytes of memory on the process' heap and returns a pointer to it that can subsequently be used by the process to hold up to `size` bytes. The contents of the memory are undefined.

void *calloc(size_t nmemb, size_t size)

calloc allocates `nmemb*size` bytes of memory on the process' heap and returns a pointer to it that can subsequently be used by the process to hold up to `size` bytes. The contents of the memory are set to zero.

void *realloc(void *ptr, size_t size)

realloc changes the size of the memory block pointed to by `ptr` to `size` bytes. The contents are copied up to `min(size, old size)`, the rest is undefined.

void free(void *ptr)

free explicitly frees a previously allocated block of memory.

Use the Unix manual to learn more about malloc, calloc, realloc, and free:

```
~/linklab $ man malloc
```

Dynamic memory management functions are defined in the standard library. Include `stdlib.h` in your program to have access to `malloc`, `calloc`, `realloc`, and `free`.

Example

The following example program demonstrates how dynamic memory management functions can be used.

```
#include <stdlib.h>

void main(void) {
    void *p;
    char *str;
    int *A;

    // allocated 1024 bytes of memory
    p = malloc(1024);

    // allocated an integer array with 500 integer
    A = (int*)calloc(500, sizeof(int));

    // allocate a string with 16 characters...
    str = (char*)malloc(16*sizeof(char));

    // ...then resize that string to hold 512 characters
    str = (char*)realloc(str, 512*sizeof(char));

    // finally, free all allocated memory
    free(p);
    free(A);
    free(str);
}
```

example1.c

Note that all allocators return an untyped pointer (`void*`) that needs to be converted to the correct type in order to prevent compiler warnings.

Part 1: Tracing Dynamic Memory Allocation (20 Points)

Subdirectory: part1/

In this part, your job is to trace all dynamic memory allocations/deallocations of a program. Print out the name, the arguments, and the return value of each call to `malloc`, `calloc`, `realloc`, and `free`. When the program ends, print statistics about memory allocation (number of bytes allocated over the course of the entire program, average size of allocation). You can ignore freed bytes when calling `realloc` (just add all allocated bytes).

For the program `test1.c` given below

```
#include <stdlib.h>

void main(void) {
    void *a;

    a = malloc(1024);
    a = malloc(32);
    free(malloc(1));
    free(a);
}
```

test1.c

the following output should be generated (the pointer values may differ from system to system):

```
~/linklab/part1 $ make run test1
[0001] Memory tracer started.
[0002]          (nil) : malloc( 1024 ) = 0xb87010
[0003]          (nil) : malloc( 32 ) = 0xb87420
[0004]          (nil) : malloc( 1 ) = 0xb87450
[0005]          (nil) : free( 0xb87450 )
[0006]          (nil) : free( 0xb87420 )
[0007]
[0008] Statistics
[0009]   allocated_total      1057
[0010]   allocated_avg       352
[0011]   freed_total        0
[0012]
[0013] Memory tracer stopped.
~/linklab/part1 $
```

Leave the statistics on the total number of freed memory bytes (`freed_total`) at 0 for now.

You can start from scratch or implement your solution by extending the skeleton provided in `~/linklab/part1/memtrace.c`. You do not need to modify `callinfo.c/h` for this part.

Use the logging facilities provided in `~/linklab/utils/memlog.c/h`. This will ensure that your output looks exactly as above and we can automatically test your submission for correctness.

Part 2: Tracing Unfreed Memory (20 Points)

Subdirectory: part2/

The tracer from part 1 is quite useful, but it has a serious shortcoming: it cannot check whether all allocated memory has been freed. In this part, we add this functionality. While the program is running, keep track of all allocated blocks and check which ones get deallocated. When the program ends, print a list of those blocks that were not deallocated. In addition, also compute and output statistics about the total number of freed memory bytes.

For the program `test1.c` shown on the previous page, the output should look as follows:

```
~/linklab/part2 $ make run test1
[0001] Memory tracer started.
[0002]      (nil) : malloc( 1024 ) = 0x2415060
[0003]      (nil) : malloc( 32 ) = 0x24154c0
[0004]      (nil) : malloc( 1 ) = 0x2415540
[0005]      (nil) : free( 0x2415540 )
[0006]      (nil) : free( 0x24154c0 )
[0007]
[0008] Statistics
[0009]   allocated_total      1057
[0010]   allocated_avg        352
[0011]   freed_total        33
[0012]
[0013] Non-deallocated memory blocks
[0014]   block      size      ref cnt   caller
[0015]   0x2415060   1024        1     ??? : 0
[0016]
[0017] Memory tracer stopped.
~/linklab/part2 $
```

You need to keep track of blocks to implement this functionality. We provide a memory block list in `~/linklab/utils/memlist.c/h` so that you don't have to write a linked list yourself (of course, you can write your own if you like). The documentation can be found in the header file.

Extend your tracer from part 1 by copying the source file `memtrace.c` into the directory of this part. You do not need to modify `callinfo.c/h` for this part.

Part 3: Pinpointing Call Locations (20 Points)

Subdirectory: part3/

To be truly useful, we need to know where in the program code the memory allocations/deallocations occur. Implement this functionality by printing the function name and the offset of the instruction calling the dynamic memory management function.

The logging functions already print the PC using the function `get_callinfo()` defined in `callinfo.h` and implemented in `callinfo.c/h`.

```
//  
// return the PC of the callsite to the dynamic memory management  
function  
//  
//  fname      pointer to character array to hold function name  
//  fnlen      length of character array  
//  ofs        pointer to offset to hold PC offset into function  
//  
// returns  
//  0          on success  
//  <0        on error  
//  
int get_callinfo(char *fname, size_t fnlen, unsigned long long *ofs);
```

~/linklab/part3/callinfo.h

```
int get_callinfo(char *fname, size_t fnlen,  
                 unsigned long long *ofs)  
{  
    return -1;  
}
```

~/linklab/part3/callinfo.c

Currently, the implementation is empty and returns -1 indicating error, which is why the PC is shown as "(nil)" in the previous outputs. Your job is to fill it with useful content so that it returns the name and offset of the caller to the `malloc/calloc/realloc/free` library function.

The output of test1.c should now look as follows:

```
~/linklab/part3 $ make run test1
[0001] Memory tracer started.
[0002]      main:6 : malloc( 1024 ) = 0x14f0060
[0003]      main:10 : malloc( 32 ) = 0x14f04c0
[0004]      main:1d : malloc( 1 ) = 0x14f0540
[0005]      main:25 : free( 0x14f0540 )
[0006]      main:2d : free( 0x14f04c0 )
[0007]
[0008] Statistics
[0009]   allocated_total      1057
[0010]   allocated_avg        352
[0011]   freed_total         33
[0012]
[0013] Non-deallocated memory blocks
[0014]   block                size      ref cnt    caller
[0015]   0x14f0060            1024         1      main:6
[0016]
[0017] Memory tracer stopped.
~/linklab/part3 $
```

Indeed, we see that our memory tracer now pinpoints the exact locations of the calls by looking at the disassembly of ~/linklab/test/test1

```
~/linklab/test $ objdump -d test1
...
00000000004004a0 <main>:
 4004a0: 53                      push    %rbx
 4004a1: bf 00 04 00 00         mov     $0x400,%edi
 4004a6: e8 e5 ff ff ff       callq   400490 <malloc@plt>
 4004ab: bf 20 00 00 00         mov     $0x20,%edi
 4004b0: e8 db ff ff ff       callq   400490 <malloc@plt>
 4004b5: bf 01 00 00 00         mov     $0x1,%edi
 4004ba: 48 89 c3               mov     %rax,%rbx
 4004bd: e8 ce ff ff ff       callq   400490 <malloc@plt>
 4004c2: 48 89 c7               mov     %rax,%rdi
 4004c5: e8 96 ff ff ff       callq   400460 <free@plt>
 4004ca: 48 89 df               mov     %rbx,%rdi
 4004cd: e8 8e ff ff ff       callq   400460 <free@plt>
 4004d2: 31 c0                 xor     %eax,%eax
 4004d4: 5b                      pop     %rbx
 4004d5: c3                      retq
 4004d6: 66 2e 0f 1f 84 00 00   nopw    %cs:0x0(%rax,%rax,1)
 4004dd: 00 00 00
~/linklab/test $
```

Bonus: Detect and Ignore Illegal Deallocations (+10 Points)

Subdirectory: bonus/

Note: to work on the bonus assignment, you have to complete part 2 first.

Some programs call free more than once on a memory block.

```
#include <stdlib.h>

void main(void) {
    void *a;

    a = malloc(1024);
    free(a);
    free(a);
    free((void*)0x1706e90);
}
```

test4.c

This programming error results in the process being aborted as shown below:

```
~/linklab/test $ ./test4
*** Error in `./test4': double free or corruption (top): 0x000000025da010 ***
===== Backtrace: =====
/lib64/libc.so.6(+0x72603) [0x7f09c84f4603]
/lib64/libc.so.6(+0x77ee6) [0x7f09c84f9ee6]
...
Aborted
~/linklab/test $
```

Extend your tracer from part 2 by detecting and ignoring deallocations of not allocated memory blocks. For the program test4.c shown above, the output should look as follows:

```
~/linklab/bonus $ make run test4
[0001] Memory tracer started.
[0002]      main:6 : malloc( 1024 ) = 0x1b30060
[0003]      main:11 : free( 0x1b30060 )
[0004]      main:19 : free( 0x1b30060 )
[0005]      *** DOUBLE_FREE *** (ignoring)
[0006]      main:23 : free( 0x1706e90 )
[0007]      *** ILLEGAL_FREE *** (ignoring)
[0008]
[0009] Statistics
[0010]   allocated_total      1024
[0011]   allocated_avg       1024
[0012]   freed_total        1024
[0013]
[0014] Memory tracer stopped.
~/linklab/bonus $
```

The output above was generated with a completed part 3. You can work on the bonus question without solving part 3 first; in that case the output will show (nil) for the program locations.